

Experimenting with Distributed Data Processing in Haskell

Utku Demir

YOW! Lambda Jam 2019

Distributed Data Processing

- ▶ Lots of data ($>$ TB)
- ▶ Complex, flexible queries

Resilient Distributed Datasets (Apache Spark)

- ▶ **Dataset** is a multiset of rows.
- ▶ A Dataset is stored as many **Partition**'s.
- ▶ Partitions can be processed in parallel on many **Executor**'s.
- ▶ High-level combinators for transforming a Dataset
 - ▶ map, filter, groupBy, reduce eg.

- ▶ Written in Haskell.
- ▶ Borrows ideas from Apache Spark.
- ▶ Composes nicely with the existing Haskell ecosystem.
- ▶ Can be used with “Function as a Service” / “Serverless” offerings.

-XStaticPointers

- ▶ `GHC.StaticPtr`

*The language extension -XStaticPointers adds a new syntactic form 'static e', which stands for a reference to the closed expression 'e'. This reference remains valid across different processes on possibly different machines.*¹

- ▶ `distributed-closure`

*This package implements a serializable closure abstraction on top of static pointers.*²

¹GHC User Guide

²<https://github.com/tweag/distributed-closure>

In detail

```
mkPartition :: Closure (ConduitT () a (ResourceT IO) ())  
             -> Partition a
```

```
data Dataset a where  
  DExternal  :: StaticSerialise a => [Partition a] -> Dataset a
```

```
DPipe       :: (StaticSerialise a, StaticSerialise b)  
             => Closure (ConduitT a b (ResourceT IO) ())  
             -> Dataset a -> Dataset b
```

```
DPartition :: (StaticHashable k, StaticSerialise a)  
            => Int  
            -> Closure (a -> k)  
            -> Dataset a -> Dataset a
```

Aggregations

```
-- Aggregates many a's to a single b.
data Aggr a b = ...

instance StaticApply (Aggr a)
instance StaticProfunctor Aggr

aggrFromMonoid :: StaticSerialise a
               => Closure (Dict (Monoid a))
               -> Aggr a a

dAggr :: Aggr a b -> Dataset a -> IO b
dGroupedAggr :: StaticHashable k
              => (a -> k) -> Aggr a b
              -> Dataset (k, b)
```

- ▶ LocalProcessBackend
- ▶ distributed-dataset-aws
 - ▶ Uses AWS Lambda to run executors and S3 to exchange information.
 - ▶ Scales well, cost-effective
 - ▶ No infrastructure necessary

Example

```
ghArchive (fromGregorian 2018 1 1, fromGregorian 2018 12 31)
  & dConcatMap (static (\e ->
    let author = e ^. gheActor . ghaLogin
        commits = e ^.. gheType . _GHPushEvent
            . ghpepCommits . traverse . ghcMessage
    in map (author, ) commits
  ))
  & dFilter (static (\(_, commit) ->
    "cabal" `T.isInfixOf` T.toLower commit
  ))
  & dGroupedAggr 50 (static fst) dCount
  & dAggr (dTopK (static Dict) 20 (static snd))
  >>= mapM_ (liftIO . print)
```

- ▶ 126 GB compressed, 909 GB uncompressed JSON
- ▶ 2190 executors
- ▶ ~ 2 minutes, including the time to deploy infrastructure

Future

- ▶ Performance optimisations
- ▶ More backends
 - ▶ static machines, Apache OpenWhisk, YARN
- ▶ Documentation, tutorials
- ▶ ML algorithms
 - ▶ clustering, recommendation

Alternatives in Haskell (as far as I know)

Sparkle

A library for writing resilient analytics applications in Haskell that scale to thousands of nodes, using Spark and the rest of the Apache ecosystem under the hood.

Hadron

Construct and run Hadoop MapReduce programs in Haskell

HSpark

A port of Apache Spark to Haskell using distributed process

Thanks!

- ▶ <https://github.com/utdemir/distributed-dataset>
- ▶ me@utdemir.com

Questions?

An external data source

```
import Network.HTTP.Simple
import Data.Conduit.Zlib (ungzip)
import Data.Conduit.JSON.NewlineDelimited as NDJ

data GHEvent = ... deriving FromJSON

urlToPartition :: String -> Partition GHEvent
urlToPartition url' = mkPartition $(\url -> do
  req <- parseRequest url
  httpSource req getResponseBody
    .| ungzip
    .| NDJ.eitherParser @_ @GHEvent
    .| C.mapM (either fail return)
) `cap` cpure (static Dict) url'
```

How to aggregate

```
input & groupedAggr 3 (static getColor) (dSum (static Dict))
```

- ▶ Input

- ▶ Partition 1: [3, 5, 2, 1]
- ▶ Partition 2: [3, 7, 2]
- ▶ Partition 3: [1, 2, 8]

- ▶ Aggregation Step 1:

- ▶ Partition 1: [5, 5, 1]
- ▶ Partition 2: [5, 7]
- ▶ Partition 3: [1, 2, 8]

- ▶ Shuffle!

- ▶ Partition 1: [5, 7, 2]
- ▶ Partition 2: [5, 1, 1]
- ▶ Partition 3: [5, 8]

- ▶ Aggregation Step 2:

- ▶ Partition 1: [14]
- ▶ Partition 2: [6, 1]
- ▶ Partition 3: [13]

Composing Aggr's

```
dConstAggr :: (Typeable a, Typeable t)
            => Closure a -> Aggr t a
```

```
dSum :: StaticSerialise a
      => Closure (Dict (Num a)) -> Aggr a a
```

```
dCount :: Typeable a => Aggr a Integer
dCount = static (const 1) `staticLmap` dSum (static Dict)
```

```
dAvg :: Aggr Double Double
dAvg = dConstAggr (static (/))
      `staticApply` dSum (static Dict)
      `staticApply` staticMap (static realToFrac) dCount
```


Results

```
("peti",899)
("haskell-pushbot",535)
("bgamari",418)
("phadej",307)
("23Skidoo",208)
("alanz",174)
("edolstra",141)
("quasicomputational",136)
("jneira",135)
("hvr",133)
("Ericson2314",133)
("felixonmars-bot",130)
("DanielG",129)
("philderbeast",120)
("coreyoconnor",115)
("rcaballeromx",109)
("...",100)
```